

CompPerf (/github/AaronTHolt/CompPerf/tree/master)

/ ANOVA-and-DOE-assignment.ipynb (/github/AaronTHolt/CompPerf/tree/master/ANOVA-and-DOE-assignment.ipynb)

ANOVA and Design of Experiments

Problem #1 - DoE

Use the following data to conduct a Design of Experiments analysis by hand (i.e. not using R). This is a 2^3 design, which is about the limit of complexity you'd want to handle by hand. You should produce a table similar to those in page 52 the DOE slides. I've provided a generic 2^3 table below; you would augment it to include the data below.

```
In [158]: temporaryFile <- tempfile()
download.file("https://www.dropbox.com/s/avt4011oc9f8qk6/3-factor-sign-table.csv?dl=0",
              destfile=temporaryFile, method="curl", extra="-L")
signTable = read.csv(temporaryFile)
```

```
In [159]: signTable$Y = c(100, 15, 40, 30, 120, 110, 20, 50)
signTable
```

Out[159]:

| | I | A | B | C | AB | AC | BC | ABC | Y |
|---|---|----|----|----|----|----|----|-----|-----|
| 1 | 1 | -1 | -1 | -1 | 1 | 1 | 1 | -1 | 100 |
| 2 | 1 | -1 | -1 | 1 | 1 | -1 | -1 | 1 | 15 |
| 3 | 1 | -1 | 1 | -1 | -1 | 1 | -1 | 1 | 40 |
| 4 | 1 | -1 | 1 | 1 | -1 | -1 | 1 | -1 | 30 |
| 5 | 1 | 1 | -1 | -1 | -1 | 1 | 1 | 1 | 120 |
| 6 | 1 | 1 | -1 | 1 | -1 | 1 | -1 | -1 | 110 |
| 7 | 1 | 1 | 1 | -1 | 1 | -1 | -1 | -1 | 20 |
| 8 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 50 |

You should compute terms q_A , q_B , ..., q_{ABC} , or the effects size of each term.

```
In [160]: qA = (-100-15-40-30+120+110+20+50)/8
qB = (-100-15+40+30-120-110+20+50)/8
qC = (-100+15-40+30-120+110-20+50)/8
qAB = (+100+15-40-30-120-110+20+50)/8
qAC = (+100-15+40-30-120+110-20+50)/8
qBC = (+100-15-40+30+120-110-20+50)/8
qABC = (-100+15+40-30+120-110-20+50)/8

effects = c(qA, qB, qC, qAB, qAC, qBC, qABC)
effects
```

```
Out[160]: 14.375 -25.625 -9.375 -14.375 14.375 14.375 -4.375
```

Now, compute the percentage variation due to each

```
In [161]: #Find sst
eff = c(qA,qB,qC,qAB,qAC,qBC,qABC)
eff2 = eff^2
(SST = sum(8*eff2))

#Find percentage variation due to each
(8*eff2/SST)
(8*eff2)
```

Out[161]: 12721.875

```
Out[161]: 0.129943502824859 0.41292065831491 0.0552689756816507 0.129943502824859 0.129943502824859
0.129943502824859 0.0120363547040039
```

```
Out[161]: 1653.125 5253.125 703.125 1653.125 1653.125 1653.125 153.125
```

You can now compare it to the results using the **lm** and **anova** functions

```
In [162]: #linear model
model = lm(Y ~ A + B + C + AB + AC + BC + ABC , data=signTable)
(model)

#anova using aov
fit <- aov(Y ~ A*B*C, data=signTable)
(fit)

#anova using anova
fit2 <- anova(model)
(fit2)
```

```
Out[162]: Call:
lm(formula = Y ~ A + B + C + AB + AC + BC + ABC, data = signTable)
```

```
Coefficients:
(Intercept)          A           B           C          AB          AC
      60.625      14.375     -25.625     -9.375     -14.375      14.375
          BC         ABC
      14.375      -4.375
```

```
Out[162]: Call:
aov(formula = Y ~ A * B * C, data = signTable)
```

```
Terms:
          A           B           C          A:B          A:C          B:C          A:B:C
Sum of Squares 1653.125 5253.125 703.125 1653.125 1653.125 1653.125 153.125
Deg. of Freedom      1         1         1         1         1         1         1
```

Estimated effects are balanced

Warning message:
In aov.lm(model): ANOVA F-tests on an essentially perfect fit are unreliable

```
Out[162]:
```

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|------------------|----|----------|----------|---------|--------|
| A | 1 | 1653.125 | 1653.125 | NaN | NaN |
| B | 1 | 5253.125 | 5253.125 | NaN | NaN |
| C | 1 | 703.125 | 703.125 | NaN | NaN |
| AB | 1 | 1653.125 | 1653.125 | NaN | NaN |
| AC | 1 | 1653.125 | 1653.125 | NaN | NaN |
| BC | 1 | 1653.125 | 1653.125 | NaN | NaN |
| ABC | 1 | 153.125 | 153.125 | NaN | NaN |
| Residuals | 0 | 0 | NaN | NA | NA |

Problem #2 - 2-Level Dhrystone

This file (dhry-2level.csv) contains the following table

```
In [163]: temporaryFile <- tempfile()
download.file("https://www.dropbox.com/s/leabtdr5tsrtlaf/dhry-2level.csv?dl=0",
             destfile=temporaryFile, method="curl", extra="-L")
dhry2lvl = read.csv(temporaryFile)
#dhry2lvl # - print this out and it will have 24 rows
```

This is the result of running a 3×2^3 experimental design for evaluating the importance of certain compiler optimizations -- you can ignore their meaning for this problem (this is a subset of data from an experiment described below). Use R to conduct an analysis of variance for this data.

a) Set up a linear model. Treat the a, f and g level as categorical factors. You can do this using e.g.

```
data$bits = factor(data$bits).
```

```
In [164]: #make them catigorical factors
dhry2lvl$a = factor(dhry2lvl$a)
dhry2lvl$f = factor(dhry2lvl$f)
dhry2lvl$g = factor(dhry2lvl$g)

#linear model
m = lm(mips ~ a*f*g, data=dhry2lvl)
```

b) Report the linear model.

```
In [165]: #report linear model
(m)
summary(m)
```

```
Out[165]: Call:
lm(formula = mips ~ a * f * g, data = dhry2lvl)

Coefficients:
(Intercept)          a1           f1           g1          a1:f1          a1:g1
    1440.5070      4.1047   -114.7363     4.8220     0.6523    17.5640
          f1:g1      a1:f1:g1
         8.5200     -34.2030
```

```
Out[165]: Call:
lm(formula = mips ~ a * f * g, data = dhry2lvl)

Residuals:
    Min       1Q   Median       3Q      Max
-15.319  -6.148   1.662   5.069  17.420

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  1440.5070    6.3922  225.355 < 2e-16 ***
a1             4.1047     9.0399   0.454  0.6559
f1            -114.7363    9.0399 -12.692 9.09e-10 ***
g1             4.8220     9.0399   0.533  0.6011
a1:f1          0.6523    12.7843   0.051  0.9599
a1:g1         17.5640    12.7843   1.374  0.1884
f1:g1          8.5200    12.7843   0.666  0.5146
a1:f1:g1     -34.2030    18.0798  -1.892  0.0768 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 11.07 on 16 degrees of freedom
Multiple R-squared:  0.9777,    Adjusted R-squared:  0.968
F-statistic: 100.4 on 7 and 16 DF,  p-value: 5.093e-12
```

c) Note that R reports this differently for categorical data. The report should be missing terms for "a", "f" and so on. Why is that? What does the intercept represent? Note that categorical values must take one some value and the intercept needs to include the results of some value(s).

In [166]:

```
print("The a, f terms are split into a linear and constant part.")
print("The constants of a, f... are absorbed in the intercept.")
print("The intercept of the lm is the the every factors constant term combined.")

[1] "The a, f terms are split into a linear and constant part."
[1] "The constants of a, f... are absorbed in the intercept."
[1] "The intercept of the lm is the the every factors constant term combined."
```

d) Determine the percentage of variation attributable to each factor. You can do this using the anova table as described in the lecture notes.

In [167]:

```
#anova
anova(m)
anova_m1 <- anova(m)

#pct variation
sst = sum(anova(m)$Sum)
pct_var = anova(m)$Sum/sst
#rename columns
anova_m1 <- setNames(cbind(rownames(anova_m1), anova_m1, row.names = NULL),
  c("factor", "Df", "SumSq", "MeanSq", "FValue", "PR(>F)"))

#put results in table
ff <- data.frame(matrix(pct_var, nrow=1, ncol=8))
colnames(ff) <- anova_m1$factor
attr(ff,"title") <- "Percentage Variation"

#display table
print("Percentage Variation of residuals")
(ff)
```

Out[167]:

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|------------------|----|-----------|-----------|-------------|--------------|
| a | 1 | 130.4101 | 130.4101 | 1.063884 | 0.3176712 |
| f | 1 | 84539.45 | 84539.45 | 689.6719 | 1.385664e-14 |
| g | 1 | 520.4198 | 520.4198 | 4.245579 | 0.05599911 |
| a:f | 1 | 405.8626 | 405.8626 | 3.311023 | 0.08757807 |
| a:g | 1 | 0.3208594 | 0.3208594 | 0.002617567 | 0.9598296 |
| f:g | 1 | 110.4632 | 110.4632 | 0.9011577 | 0.356591 |
| a:f:g | 1 | 438.692 | 438.692 | 3.578844 | 0.07676194 |
| Residuals | 16 | 1961.268 | 122.5792 | NA | NA |

```
[1] "Percentage Variation of residuals"
```

Out[167]:

| | a | f | g | a:f | a:g | f:g | a:f:g | Residuals |
|----------|-------------|-----------|-------------|------------|--------------|-------------|--------------|------------------|
| 1 | 0.001480136 | 0.9595101 | 0.005906687 | 0.00460648 | 3.641706e-06 | 0.001253741 | 0.004979088 | 0.0222601 |

e) Compute the 95% confidence interval for the "opt" factor. The "anova" function doesn't do this, but you can see something similar using

```
confint( aov (l ) )
```

where 'l' is your linear model. For models involving factors, the confidence interval is expressed for the different levels in each categorical factor.

In [168]:

```
#find confidence interval  
confint(aov(m))  
print("There is no opt factor so here is the confidence interval for all of them.")
```

Out[168]:

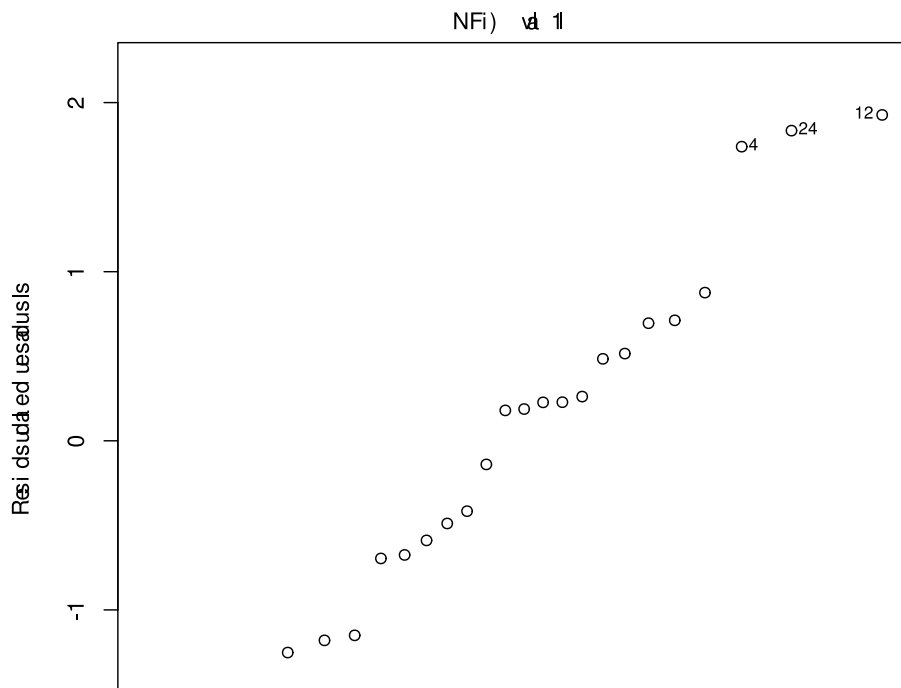

| | 2.5 % | 97.5 % |
|--------------------|------------|-----------|
| (Intercept) | 1426.956 | 1454.058 |
| a1 | -15.05903 | 23.26836 |
| f1 | -133.90003 | -95.57264 |
| g1 | -14.3417 | 23.9857 |
| a1:f1 | -26.44923 | 27.75389 |
| a1:g1 | -9.537558 | 44.665558 |
| f1:g1 | -18.58156 | 35.62156 |
| a1:f1:g1 | -72.530392 | 4.124392 |


```
[1] "There is no opt factor so here is the confidence interval for all of them."
```


f) Plot out the diagnostic plots for the model. You can do this using "plot(aov(l))" where "l" is your model. Can you justify that your model is reasonable?

In [169]:

```
#plot anova stuff
plot(aov(m))
print("The residuals are fairly small and don't seem to have a pattern associated with them.")
print("Also they are evenly distributed in the + and - side.")
print("Furthermore the QQ plot is pretty linear.")
print("These all point to a reasonable model.")
print("Only worrisome point is the lack of data.")
```







Problem #3 -

The file OUTPUT-optlevel (see below) contains data from a 422 study with 2 replicates. The first column is the trial number (you shouldn't include this in your model, or you won't get any "residuals"). The first factor ('opt') has four levels (-O0, -O1, -O2, -O3). The second factor ('bits') has two levels (specifying -m32 or -m64 for compilation as a 32-bit or 64-bit application) and the third ('benchmark') has two levels (dhry11 or dhry21). In each case, the dependent variable is the performance in VAX MIPS (this is the output of the Dhrystone program). Use R to conduct an analysis of variance for this data.

```
In [170]: temporaryFile <- tempfile()
download.file("https://www.dropbox.com/s/l0lenc8prwfo7e3/OUTPUT-optlevels.orig?dl=0",
             destfile=temporaryFile, method="curl", extra="-L")
optlevels = read.csv(temporaryFile)
#optlevels # - this will have 32 rows
```

a) Set up a linear model. Treat the bits, prog and optimization level as categorical factors.

```
In [171]: #make them catigorical
optlevels$opt = factor(optlevels$opt)
optlevels$bits = factor(optlevels$bits)
optlevels$prog = factor(optlevels$prog)

#linear model
m2 = lm(mips ~ opt*bits*prog, data=optlevels)
```

b) Report the linear model. Note that R reports this differently for categorical data. The report should be missing terms for "opt0", "bits0" and so on. Why is that? What does the intercept represent?


```
In [172]: #diplay lm
(m2)
summary(m2)
print("The opt0 and bits0 terms are absorbed into the intercept.")
print("The intercept of the lm is the the every factors constant term combined.")
```

```
Out[172]: Call:
lm(formula = mips ~ opt * bits * prog, data = optlevels)

Coefficients:
(Intercept)          opt1          opt2          opt3
      3077.7          504.3          4911.3          4486.1
      bits1          prog1      opt1:bits1      opt2:bits1
      2329.9          462.3          -634.6          -808.9
      opt3:bits1      opt1:prog1      opt2:prog1      opt3:prog1
      -187.5          -375.8          -848.0          -194.8
      bits1:prog1      opt1:bits1:prog1      opt2:bits1:prog1      opt3:bits1:prog1
      -472.5          383.9          1111.2          204.2
```

```
Out[172]: Call:
lm(formula = mips ~ opt * bits * prog, data = optlevels)
```

```
Residuals:
      Min       1Q   Median       3Q      Max
-278.14  -71.26    0.00   71.26  278.14
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    3077.7      134.3   22.921 1.16e-13 ***
opt1             504.3      189.9    2.656 0.01726 *
opt2            4911.3      189.9   25.864 1.76e-14 ***
opt3            4486.1      189.9   23.624 7.24e-14 ***
bits1           2329.9      189.9   12.270 1.49e-09 ***
prog1            462.3      189.9    2.435 0.02700 *
opt1:bits1      -634.6      268.5   -2.363 0.03113 *
opt2:bits1      -808.9      268.5   -3.012 0.00827 **
opt3:bits1      -187.5      268.5   -0.698 0.49517
opt1:prog1      -375.8      268.5   -1.400 0.18074
opt2:prog1      -848.0      268.5   -3.158 0.00609 **
opt3:prog1      -194.8      268.5   -0.725 0.47867
bits1:prog1     -472.5      268.5   -1.760 0.09759 .
opt1:bits1:prog1  383.9      379.8    1.011 0.32720
opt2:bits1:prog1 1111.2      379.8    2.926 0.00989 **
opt3:bits1:prog1  204.2      379.8    0.538 0.59822
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 189.9 on 16 degrees of freedom
Multiple R-squared:  0.9968,    Adjusted R-squared:  0.9937
F-statistic: 328.9 on 15 and 16 DF,  p-value: < 2.2e-16
```

```
[1] "The opt0 and bits0 terms are absorbed into the intercept."
[1] "The intercept of the lm is the the every factors constant term combined."
```

c) Determine the percentage of variation attributable to each factor. You can do this using the anova table.

In [173]:

```
#anova
anova(m2)
anova_m2 <- anova(m2)

#pct variation
sst2 = sum(anova(m2)$Sum)
pct_var2 = anova(m2)$Sum/sst2
#(pct_var2)

#change col names
anova_m2 <- setNames(cbind(rownames(anova_m2), anova_m2, row.names = NULL),
  c("factor", "Df", "SumSq", "MeanSq", "FValue", "PR(>F)"))

#put results in table
ff <- data.frame(matrix(pct_var2, nrow=1, ncol=8))
colnames(ff) <- anova_m2$factor
attr(ff,"title") <- "Percentage Variation"

#display table
print("Percentage Variation of residuals")
(ff)
```

Out[173]:

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|---------------|----|-----------|----------|-----------|--------------|
| opt | 3 | 148353389 | 49451130 | 1371.402 | 1.690382e-19 |
| bits | 1 | 28829066 | 28829066 | 799.5012 | 4.349935e-15 |
| prog | 1 | 56143.58 | 56143.58 | 1.557 | 0.2300588 |
| opt:bits | 3 | 229529.3 | 76509.76 | 2.121804 | 0.1376009 |
| opt:prog | 3 | 93947.14 | 31315.71 | 0.868462 | 0.4777577 |
| bits:prog | 1 | 4552.321 | 4552.321 | 0.1262471 | 0.7269961 |
| opt:bits:prog | 3 | 350977.3 | 116992.4 | 3.244489 | 0.04975912 |
| Residuals | 16 | 576941.1 | 36058.82 | NA | NA |

[1] "Percentage Variation of residuals"

Out[173]:

| | opt | bits | prog | opt:bits | opt:prog | bits:prog | opt:bits:prog | Residuals |
|---|-----------|-----------|--------------|-------------|--------------|--------------|---------------|-------------|
| 1 | 0.8311368 | 0.1615123 | 0.0003145394 | 0.001285918 | 0.0005263306 | 2.550398e-05 | 0.001966319 | 0.003232262 |

d) Compute the 95% confidence interval for the "opt" factor. The "anova" function doesn't do this, but you can see something similar using

```
confint( aov ( l ) )
```

where 'l' is your linear model. For models involving factors, the confidence interval is expressed for the different levels in each categorical factor.

In [174]:

```
#display confidence intervals
confint(aov(m2))
```


Out[174]:

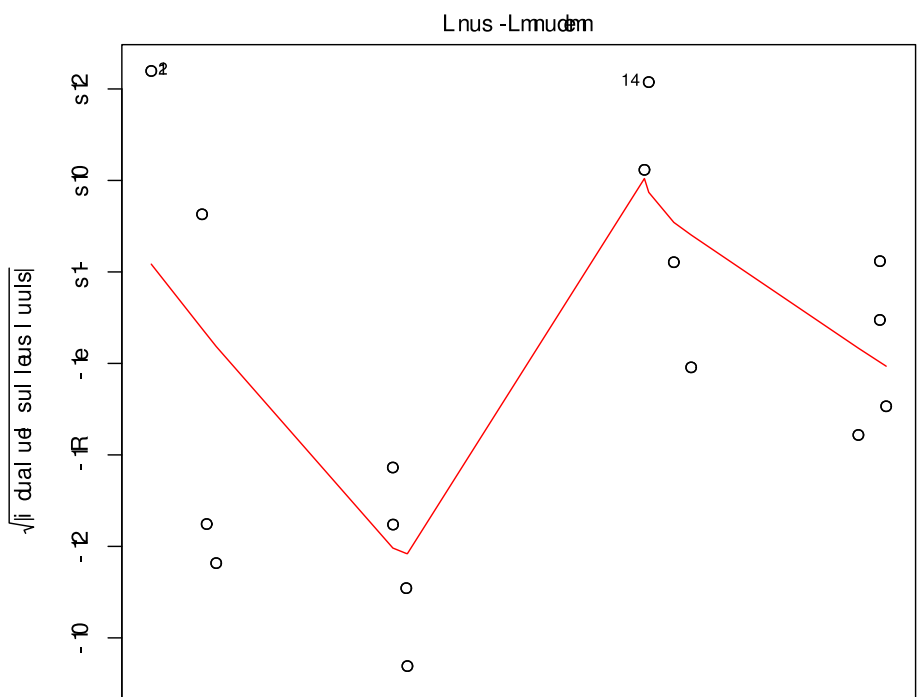

| | 2.5 % | 97.5 % |
|-------------------------|-------------|-----------|
| (Intercept) | 2793.061 | 3362.355 |
| opt1 | 101.7688 | 906.8732 |
| opt2 | 4508.722 | 5313.827 |
| opt3 | 4083.535 | 4888.639 |
| bits1 | 1927.361 | 2732.465 |
| prog1 | 59.74181 | 864.84619 |
| opt1:bits1 | -1203.84477 | -65.25523 |
| opt2:bits1 | -1378.2248 | -239.6352 |
| opt3:bits1 | -756.7533 | 381.8363 |
| opt1:prog1 | -945.1343 | 193.4553 |
| opt2:prog1 | -1417.3108 | -278.7212 |
| opt3:prog1 | -764.1043 | 374.4853 |
| bits1:prog1 | -1041.81577 | 96.77377 |
| opt1:bits1:prog1 | -421.2534 | 1188.9554 |
| opt2:bits1:prog1 | 306.1021 | 1916.3109 |
| opt3:bits1:prog1 | -600.9144 | 1009.2944 |


e) Plot out the diagnostic plots for the model. You can do this using "plot(aov(l))" where "l" is your model. Can you justify that your model is reasonable

In [175]:


```
#plot anova stuff
plot(aov(m2))
print("The residuals are fairly small and don't seem to have a pattern associated with them.")
print("Also they are evenly distributed in the + and - side.")
print("Furthermore the QQ plot is pretty linear.")
print("Worrisome point is the lack of data.")
print("Also different points have very different residuals associated with them.")
print("still seems to be a reasonable model.")
```







- [1] "The residuals are fairly small and don't seem to have a pattern associated with them."
- [1] "Also they are evenly distributed in the + and - side."
- [1] "Furthermore the QQ plot is pretty linear."
- [1] "Worrisome point is the lack of data."
- [1] "Also different points have very different residuals associated with them."
- [1] "still seems to be a reasonable model."



Problem #4 -

This file (OUTPUT-21.csv) contains a $3 * 2^7$ experiment. The data was generated from the dhrystone 2.1 benchmark on a specific machine. The factors in the experiment correspond to different optimizations as indicated in the RUN script:

```
DATA = (
  ('-fstrength-reduce', '-fno-strength-reduce'), # a
  ('-fgcse', '-fno-gcse'), # b
  ('-floop-optimize', '-fno-loop-optimize'), # c
  ('-fpeephole', '-fno-peephole'), # d
  ('-finline-functions', '-fno-inline-functions'), # e
  ('-fomit-frame-pointer', '-fno-omit-frame-pointer'), # f
  ('-fwhole-program', '-fno-whole-program') # g
)
```

each factor has two levels - either the optimization is enabled or not. In my model, I called the factors o.a, o.b, o.c, etc for the different possible optimizations.

In [176]:

```
temporaryFile <- tempfile()
download.file("https://www.dropbox.com/s/0hq1r63ym1h904q/OUTPUT-21.csv?dl=0",
              destfile=temporaryFile, method="curl", extra="-L")
fullOpt = read.csv(temporaryFile)
#fullOpt # - if you print this out, it will have 384 rows
```

a) There are several terms that contribute more variation to the results than the error/residual term, but most of the terms contribute little variation to the overall performance. Determine which single factor contributes in some meaningful way to variation in performance and report the percentage of variation explained by that single term. Explain how you did this.

In [177]:

```

#Turn into categorical
full0pt$o.a = factor(full0pt$o.a)
full0pt$o.b = factor(full0pt$o.b)
full0pt$o.c = factor(full0pt$o.c)
full0pt$o.d = factor(full0pt$o.d)
full0pt$o.e = factor(full0pt$o.e)
full0pt$o.f = factor(full0pt$o.f)
full0pt$o.g = factor(full0pt$o.g)

#linear model
m4 = lm(mips ~ o.a*o.b*o.c*o.d*o.e*o.f*o.g, data=full0pt)

#anova
anova_m4 = anova(m4)

#percentage variation
sst4 = sum(anova(m4)$Sum)
pct_var4 = anova(m4)$Sum/sst4

#find highest single factor and associated row
max(pct_var4)
max_row <- which.max(pct_var4)

#rename column names
anova_m4 <- setNames(cbind(rownames(anova_m4), anova_m4, row.names = NULL),
  c("factor", "Df", "SumSq", "MeanSq", "FValue", "PR(>F)"))

#print row
(anova_m4[max_row,])

print('The most important factor is o.f with a 64.99 percentage variation')
print("I found this by turning the factors into categorical variables,")
print("creating a linear model, performing anova analysis.")
print("Finally i calculated the percentage variation of all factors,")
print("and picked the highest one.")

```

Out[177]:

0.649967839195917

Out[177]:

| | factor | Df | SumSq | MeanSq | FValue | PR(>F) |
|---|--------|----|---------|---------|----------|--------------|
| 6 | o.f | 1 | 2077012 | 2077012 | 726.4939 | 9.973427e-77 |

```

[1] "The most important factor is o.f with a 64.99 percentage variation"
[1] "I found this by turning the factors into categorical variables,"
[1] "creating a linear model, performing anova analysis."
[1] "Finally i calculated the percentage variation of all factors,"
[1] "and picked the highest one."

```

b) What 2-factor interaction of factors contributes in some meaningful way to variation in performance. Explain how you determine which factors you're including.


```

In [178]: #Pick out 2-factor interactions
#anova_m4$factor
start_2s <- which(grepl("o.a:o.b", anova_m4$factor))[1]
end_2s <- which(grepl("o.f:o.g", anova_m4$factor))[1]

#Find highest
pct_var_2s <- pct_var4[start_2s:end_2s]
print("pct variation of most influential")
max(pct_var_2s)
print("row number of that factor")
which.max(pct_var_2s)

#get row using value, print row
max_row <- which(pct_var4 == max(pct_var_2s))
print("factors row from anova")
(anova_m4[max_row,])

#Find second highest, first zero out previous highest
pct_var_2s[13] <- 0
print("pct variation of second most influential")
max(pct_var_2s)
print("row number of that factor")
which.max(pct_var_2s)

#get row using value, print row
max_row <- which(pct_var4 == max(pct_var_2s))
print("factors row from anova")
(anova_m4[max_row,])

print("2-Factor interactions that have highest percentage variation are o.c:o.f at 0.56%,"
print("and o.a:o.b at 0.26%.")
print("I included all factors and then picked the ones with highest percentage variation")
print("In the 2-factor category.")

[1] "pct variation of most influential"

Out[178]: 0.00569151029703818

[1] "row number of that factor"

Out[178]: 13

[1] "factors row from anova"

Out[178]:


|    | factor  | Df | SumSq    | MeanSq   | FValue   | PR(>F)     |
|----|---------|----|----------|----------|----------|------------|
| 20 | o.c:o.f | 1  | 18187.57 | 18187.57 | 6.361619 | 0.01226866 |



[1] "pct variation of second most influential"

Out[178]: 0.00265805714650366

[1] "row number of that factor"

Out[178]: 1

[1] "factors row from anova"

Out[178]:


|   | factor  | Df | SumSq    | MeanSq   | FValue   | PR(>F)     |
|---|---------|----|----------|----------|----------|------------|
| 8 | o.a:o.b | 1  | 8493.985 | 8493.985 | 2.971012 | 0.08597612 |



[1] "2-Factor interactions that have highest percentage variation are o.c:o.f at 0.56%,"
[1] "and o.a:o.b at 0.26%."
[1] "I included all factors and then picked the ones with highest percentage variation"
[1] "In the 2-factor category."

```

c) There are two significant interactions that involve more than 2 factors. List those & explain how you did this.

```
In [179]: #Pick out 3-factor interactions
#anova_m4$factor
start_3s <- which(grepl("o.a:o.b:o.c", anova_m4$factor))[1]
end_3s <- which(grepl("o.e:o.f:o.g", anova_m4$factor))[1]

#Find highest
pct_var_3s <- pct_var4[start_3s:end_3s]
print("pct variation of most influential")
max(pct_var_3s)
print("row number of that factor")
which.max(pct_var_3s)

#get row using value, print row
max_row <- which(pct_var4 == max(pct_var_3s))
print("factors row from anova")
(anova_m4[max_row,])

#Find second highest, first zero out previous highest
pct_var_3s[5] <- 0
print("pct variation of second most influential")
max(pct_var_3s)
print("row number of that factor")
which.max(pct_var_3s)

#get row using value, print row
max_row <- which(pct_var4 == max(pct_var_3s))
print("factors row from anova")
(anova_m4[max_row,])

print("3-Factor interactions that have highest percentage variation are: ")
print("and o.a:o.b:o.e at 0.45% and o.a:o.b:o.c at 0.37%.")
print("I included all factors and then picked the ones with highest percentage variation")
print("in the 3-factor category.")
```

```
[1] "pct variation of most influential"
```

```
Out[179]: 0.00453043169162491
```

```
[1] "row number of that factor"
```

```
Out[179]: 5
```

```
[1] "factors row from anova"
```

```
Out[179]:
```

| | factor | Df | SumSq | MeanSq | FValue | PR(>F) |
|----|-------------|----|----------|----------|----------|------------|
| 33 | o.a:o.b:o.e | 1 | 14477.27 | 14477.27 | 5.063837 | 0.02527951 |

```
[1] "pct variation of second most influential"
```

```
Out[179]: 0.00373568245534494
```

```
[1] "row number of that factor"
```

```
Out[179]: 1
```

```
[1] "factors row from anova"
```

```
Out[179]:
```

| | factor | Df | SumSq | MeanSq | FValue | PR(>F) |
|----|-------------|----|---------|---------|----------|------------|
| 29 | o.a:o.b:o.c | 1 | 11937.6 | 11937.6 | 4.175515 | 0.04203591 |

```
[1] "3-Factor interactions that have highest percentage variation are: "
[1] "and o.a:o.b:o.e at 0.45% and o.a:o.b:o.c at 0.37%."
[1] "I included all factors and then picked the ones with highest percentage variation"
[1] "in the 3-factor category."
```

Problem #4 - reduced

I've taken the data from the previous question and prepared a $3 \times 2^{7-4}$ design table. You're left with a 24 row table (8 factors, repeated 3 times). This is in the following table.

```
In [183]: temporaryFile <- tempfile()
download.file("https://www.dropbox.com/s/n9urefzm9jzggh6/OUTPUT-21-reduced.csv?dl=0",
             destfile=temporaryFile, method="curl", extra="-L")
reducedOpt = read.csv(temporaryFile)
#reducedOpt # if you print this out, it will have 24 rows
```

a) Again, there should be a single factor that explains the most variation, but the percentage of variation explained by that factor should now be different. Compute the percentage for that single factor.

In [184]:

```
#Turn into categorical
reduced0pt$o.a = factor(reduced0pt$o.a)
reduced0pt$o.b = factor(reduced0pt$o.b)
reduced0pt$o.c = factor(reduced0pt$o.c)
reduced0pt$o.d = factor(reduced0pt$o.d)
reduced0pt$o.e = factor(reduced0pt$o.e)
reduced0pt$o.f = factor(reduced0pt$o.f)
reduced0pt$o.g = factor(reduced0pt$o.g)

#linear model
m5 = lm(mips ~ o.a*o.b*o.c*o.d*o.e*o.f*o.g, data=reduced0pt)
#m5 = lm(mips ~ o.a+o.b+o.c+o.d+o.e+o.f+o.g, data=reduced0pt)

#anova
anova_m5 = anova(m5)

#percentage variation
sst5 = sum(anova(m5)$Sum)
pct_var5 = anova(m5)$Sum/sst5

#find highest single factor, print row
max(pct_var5)
max_row <- which.max(pct_var5)
anova_m5 <- setNames(cbind(rownames(anova_m5), anova_m5, row.names = NULL),
  c("factor", "Df", "SumSq", "MeanSq", "FValue", "PR(>F)"))
(anova_m5[max_row,])

print("Factor o.f has the highest percentage of variation with 76.27%.")
```

Out[184]:

0.762742739498329

Out[184]:

| | factor | Df | SumSq | MeanSq | FValue | PR(>F) |
|---|--------|----|----------|----------|----------|--------------|
| 6 | o.f | 1 | 126683.4 | 126683.4 | 59.34043 | 9.023391e-07 |

[1] "Factor o.f has the highest percentage of variation with 76.27%."

b) Now, determine why that variation is different. Show that the "F" factor is confounded with the ABCDEG factor and explain why that changes the variation attributed to F.

In [182]:

```
library(car)
vif(m5)
print("Find Alias")
(alias(m5))

print("vif() couldn't be used to find variable inflation because there are aliased coefficients in the")
print("model. Thus alias() was used to look for correlated variables. Many variables were correlated. F")
print("and ABCDEG had some collinearity.")
```

Error in vif.default(m5): there are aliased coefficients in the model

```
[1] "Find Alias"
```

Out[182]:

Model :
 mips ~ o.a * o.b * o.c * o.d * o.e * o.f * o.g

Complete :

| | (Intercept) | o.a1 | o.b1 | o.c1 | o.d1 | o.e1 | o.f1 |
|---------------------|-------------|------|------|------|------|------|------|
| o.a1:o.b1 | -1/2 | 1/2 | 1/2 | 0 | 1/2 | 0 | 0 |
| o.a1:o.c1 | -1/2 | 1/2 | 0 | 1/2 | 0 | 1/2 | 0 |
| o.b1:o.c1 | -1/2 | 0 | 1/2 | 1/2 | 0 | 0 | 1/2 |
| o.a1:o.d1 | -1/2 | 1/2 | 1/2 | 0 | 1/2 | 0 | 0 |
| o.b1:o.d1 | -1/2 | 1/2 | 1/2 | 0 | 1/2 | 0 | 0 |
| o.c1:o.d1 | -1/2 | 0 | 0 | 1/2 | 1/2 | 0 | 0 |
| o.a1:o.e1 | -1/2 | 1/2 | 0 | 1/2 | 0 | 1/2 | 0 |
| o.b1:o.e1 | -1/2 | 0 | 1/2 | 0 | 0 | 1/2 | 0 |
| o.c1:o.e1 | -1/2 | 1/2 | 0 | 1/2 | 0 | 1/2 | 0 |
| o.d1:o.e1 | -1/2 | 0 | 0 | 0 | 1/2 | 1/2 | 1/2 |
| o.a1:o.f1 | -1/2 | 1/2 | 0 | 0 | 0 | 0 | 1/2 |
| o.b1:o.f1 | -1/2 | 0 | 1/2 | 1/2 | 0 | 0 | 1/2 |
| o.c1:o.f1 | -1/2 | 0 | 1/2 | 1/2 | 0 | 0 | 1/2 |
| o.d1:o.f1 | -1/2 | 0 | 0 | 0 | 1/2 | 1/2 | 1/2 |
| o.e1:o.f1 | -1/2 | 0 | 0 | 0 | 1/2 | 1/2 | 1/2 |
| o.a1:o.g1 | -1/2 | 1/2 | 0 | 0 | 0 | 0 | 1/2 |
| o.b1:o.g1 | -1/2 | 0 | 1/2 | 0 | 0 | 1/2 | 0 |
| o.c1:o.g1 | -1/2 | 0 | 0 | 1/2 | 1/2 | 0 | 0 |
| o.d1:o.g1 | -1/2 | 0 | 0 | 1/2 | 1/2 | 0 | 0 |
| o.e1:o.g1 | -1/2 | 0 | 1/2 | 0 | 0 | 1/2 | 0 |
| o.f1:o.g1 | -1/2 | 1/2 | 0 | 0 | 0 | 0 | 1/2 |
| o.a1:o.b1:o.c1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.d1 | -1/2 | 1/2 | 1/2 | 0 | 1/2 | 0 | 0 |
| o.a1:o.c1:o.d1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.d1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.e1 | -1/2 | 1/2 | 0 | 1/2 | 0 | 1/2 | 0 |
| o.b1:o.c1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.f1 | -1/2 | 0 | 1/2 | 1/2 | 0 | 0 | 1/2 |
| o.a1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.d1:o.e1:o.f1 | -1/2 | 0 | 0 | 0 | 1/2 | 1/2 | 1/2 |
| o.a1:o.b1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.d1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.d1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.d1:o.g1 | -1/2 | 0 | 0 | 1/2 | 1/2 | 0 | 0 |
| o.a1:o.e1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.e1:o.g1 | -1/2 | 0 | 1/2 | 0 | 0 | 1/2 | 0 |
| o.c1:o.e1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.d1:o.e1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.f1:o.g1 | -1/2 | 1/2 | 0 | 0 | 0 | 0 | 1/2 |
| o.b1:o.f1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.f1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.d1:o.f1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.e1:o.f1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.c1:o.d1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.c1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.d1:o.e1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.c1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.d1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.d1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.d1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.c1:o.d1:o.e1:o.f1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.c1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.b1:o.d1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.a1:o.c1:o.d1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.b1:o.c1:o.d1:o.g1 | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |

Jupyter Notebook Viewer

| | | | | | | | |
|------------------------------------|------|-----|-----|-----|-----|-----|-----|
| o.al:o.bl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.cl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.cl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.cl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.el | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.el:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.dl:o.el:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.dl:o.el:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.dl:o.el:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.cl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.el:o.fl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.el:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.cl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.bl:o.cl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl:o.cl:o.dl:o.el:o.fl:o.gl | -3/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 | 1/4 |
| o.al:o.bl | 0 | | | | | | |
| o.al:o.cl | 0 | | | | | | |
| o.bl:o.cl | 0 | | | | | | |
| o.al:o.dl | 0 | | | | | | |
| o.bl:o.dl | 0 | | | | | | |
| o.cl:o.dl | 1/2 | | | | | | |
| o.al:o.el | 0 | | | | | | |
| o.bl:o.el | 1/2 | | | | | | |
| o.cl:o.el | 0 | | | | | | |
| o.dl:o.el | 0 | | | | | | |
| o.al:o.fl | 1/2 | | | | | | |
| o.bl:o.fl | 0 | | | | | | |
| o.cl:o.fl | 0 | | | | | | |
| o.dl:o.fl | 0 | | | | | | |
| o.el:o.fl | 0 | | | | | | |
| o.al:o.gl | 1/2 | | | | | | |
| o.bl:o.gl | 1/2 | | | | | | |
| o.cl:o.gl | 1/2 | | | | | | |
| o.dl:o.gl | 1/2 | | | | | | |
| o.el:o.gl | 1/2 | | | | | | |
| o.fl:o.gl | 1/2 | | | | | | |
| o.al:o.bl:o.cl | 1/4 | | | | | | |
| o.al:o.bl:o.dl | 0 | | | | | | |
| o.al:o.cl:o.dl | 1/4 | | | | | | |
| o.bl:o.cl:o.dl | 1/4 | | | | | | |
| o.al:o.bl:o.el | 1/4 | | | | | | |
| o.al:o.cl:o.el | 0 | | | | | | |
| o.bl:o.cl:o.el | 1/4 | | | | | | |
| o.al:o.dl:o.el | 1/4 | | | | | | |
| o.bl:o.dl:o.el | 1/4 | | | | | | |
| o.cl:o.dl:o.el | 1/4 | | | | | | |
| o.al:o.bl:o.fl | 1/4 | | | | | | |
| o.al:o.cl:o.fl | 1/4 | | | | | | |
| o.bl:o.cl:o.fl | 0 | | | | | | |

| | |
|-------------------------------|-----|
| o.a1:o.d1:o.fl | 1/4 |
| o.b1:o.d1:o.fl | 1/4 |
| o.c1:o.d1:o.fl | 1/4 |
| o.a1:o.el:o.fl | 1/4 |
| o.b1:o.el:o.fl | 1/4 |
| o.c1:o.el:o.fl | 1/4 |
| o.d1:o.el:o.fl | 0 |
| o.a1:o.b1:o.gl | 1/4 |
| o.a1:o.c1:o.gl | 1/4 |
| o.b1:o.c1:o.gl | 1/4 |
| o.a1:o.d1:o.gl | 1/4 |
| o.b1:o.d1:o.gl | 1/4 |
| o.c1:o.d1:o.gl | 1/2 |
| o.a1:o.el:o.gl | 1/4 |
| o.b1:o.el:o.gl | 1/2 |
| o.c1:o.el:o.gl | 1/4 |
| o.d1:o.el:o.gl | 1/4 |
| o.a1:o.fl:o.gl | 1/2 |
| o.b1:o.fl:o.gl | 1/4 |
| o.c1:o.fl:o.gl | 1/4 |
| o.d1:o.fl:o.gl | 1/4 |
| o.el:o.fl:o.gl | 1/4 |
| o.a1:o.b1:o.c1:o.d1 | 1/4 |
| o.a1:o.b1:o.c1:o.el | 1/4 |
| o.a1:o.b1:o.d1:o.el | 1/4 |
| o.a1:o.c1:o.d1:o.el | 1/4 |
| o.b1:o.c1:o.d1:o.el | 1/4 |
| o.a1:o.b1:o.c1:o.fl | 1/4 |
| o.a1:o.b1:o.d1:o.fl | 1/4 |
| o.a1:o.c1:o.d1:o.fl | 1/4 |
| o.b1:o.c1:o.d1:o.fl | 1/4 |
| o.a1:o.b1:o.el:o.fl | 1/4 |
| o.a1:o.c1:o.el:o.fl | 1/4 |
| o.b1:o.c1:o.el:o.fl | 1/4 |
| o.a1:o.d1:o.el:o.fl | 1/4 |
| o.b1:o.d1:o.el:o.fl | 1/4 |
| o.c1:o.d1:o.el:o.fl | 1/4 |
| o.a1:o.b1:o.c1:o.gl | 1/4 |
| o.a1:o.b1:o.d1:o.gl | 1/4 |
| o.a1:o.c1:o.d1:o.gl | 1/4 |
| o.b1:o.c1:o.d1:o.gl | 1/4 |
| o.a1:o.b1:o.el:o.gl | 1/4 |
| o.a1:o.c1:o.el:o.gl | 1/4 |
| o.b1:o.c1:o.el:o.gl | 1/4 |
| o.a1:o.d1:o.el:o.gl | 1/4 |
| o.b1:o.d1:o.el:o.gl | 1/4 |
| o.c1:o.d1:o.el:o.gl | 1/4 |
| o.a1:o.b1:o.fl:o.gl | 1/4 |
| o.a1:o.c1:o.fl:o.gl | 1/4 |
| o.b1:o.c1:o.fl:o.gl | 1/4 |
| o.a1:o.d1:o.fl:o.gl | 1/4 |
| o.b1:o.d1:o.fl:o.gl | 1/4 |
| o.c1:o.d1:o.fl:o.gl | 1/4 |
| o.a1:o.el:o.fl:o.gl | 1/4 |
| o.b1:o.el:o.fl:o.gl | 1/4 |
| o.c1:o.el:o.fl:o.gl | 1/4 |
| o.d1:o.el:o.fl:o.gl | 1/4 |
| o.a1:o.b1:o.c1:o.d1:o.el | 1/4 |
| o.a1:o.b1:o.c1:o.d1:o.fl | 1/4 |
| o.a1:o.b1:o.c1:o.el:o.fl | 1/4 |
| o.a1:o.b1:o.d1:o.el:o.fl | 1/4 |
| o.a1:o.c1:o.d1:o.el:o.fl | 1/4 |
| o.b1:o.c1:o.d1:o.el:o.fl | 1/4 |
| o.a1:o.b1:o.c1:o.d1:o.gl | 1/4 |
| o.a1:o.b1:o.c1:o.el:o.gl | 1/4 |
| o.a1:o.b1:o.d1:o.el:o.gl | 1/4 |
| o.a1:o.c1:o.d1:o.el:o.gl | 1/4 |
| o.b1:o.c1:o.d1:o.el:o.gl | 1/4 |
| o.a1:o.b1:o.c1:o.fl:o.gl | 1/4 |
| o.a1:o.b1:o.d1:o.fl:o.gl | 1/4 |
| o.a1:o.c1:o.d1:o.fl:o.gl | 1/4 |
| o.b1:o.c1:o.d1:o.fl:o.gl | 1/4 |
| o.a1:o.b1:o.el:o.fl:o.gl | 1/4 |
| o.a1:o.c1:o.el:o.fl:o.gl | 1/4 |
| o.b1:o.c1:o.el:o.fl:o.gl | 1/4 |
| o.a1:o.d1:o.el:o.fl:o.gl | 1/4 |
| o.b1:o.d1:o.el:o.fl:o.gl | 1/4 |
| o.c1:o.d1:o.el:o.fl:o.gl | 1/4 |
| o.a1:o.b1:o.c1:o.d1:o.el:o.fl | 1/4 |
| o.a1:o.b1:o.c1:o.d1:o.el:o.gl | 1/4 |

4/7/2016

Jupyter Notebook Viewer

```
o.a1:o.b1:o.c1:o.d1:o.f1:o.g1    1/4
o.a1:o.b1:o.c1:o.e1:o.f1:o.g1    1/4
o.a1:o.b1:o.d1:o.e1:o.f1:o.g1    1/4
o.a1:o.c1:o.d1:o.e1:o.f1:o.g1    1/4
o.b1:o.c1:o.d1:o.e1:o.f1:o.g1    1/4
o.a1:o.b1:o.c1:o.d1:o.e1:o.f1:o.g1 1/4
```

```
[1] "vif() couldn't be used to find variable inflation because there are aliased coefficients in the
[1] "model. Thus alias() was used to look for correlated variables. Many variables were correlated.
[1] "and ABCDEG had some collinearity."
```

◀

In []: